

**Erik Hallström** Follow

Studied Engineering Physics and in Machine Learning at Royal Institute of Technology in Stockholm. Also...

Nov 18, 2016 · 2 min read

Using the LSTM API in TensorFlow (3/7)

In the [previous post](#) we modified our code to use the TensorFlow native RNN API. Now we will go about to build a modification of a RNN that called a “Recurrent Neural Network with Long short-term memory” or RNN-LSTM. This architecture was pioneered by [Jürgen Schmidhuber](#) among others. One problem with the RNN when using long time-dependencies (`truncated_backprop_length` is large) is the “[vanishing gradient problem](#)”. One way to counter this is using a state that is “protected” and “selective”. The RNN-LSTM remembers, forgets and chooses what to pass on and output depending on the current state and input.

Since this primarily is a practical tutorial I won’t go into more detail about the theory, I recommend reading [this article](#) again, continue with the “Modern RNN architectures”. After you have done that read and look at the figures on [this page](#). Notice that the last mentioned resource are using vector concatenation in their calculations.

In the previous article we didn’t have to allocate the internal weight matrix and bias, that was done by TensorFlow automatically “under the hood”. A LSTM RNN has many more “moving parts”, but by using the native API it will also be very simple.

Different state

A LSTM have a “cell state” and a “hidden state”, to account for this you need to remove `_current_state` on line 79 in the previous script and replace it with this:

```
1 _current_cell_state = np.zeros((batch_size, state_size))
2 _current_hidden_state = np.zeros((batch_size, state_size))
3
```

TensorFlow uses a data structure called `LSTMStateTuple` internally for its LSTM:s, where the first element in the tuple is the cell state, and

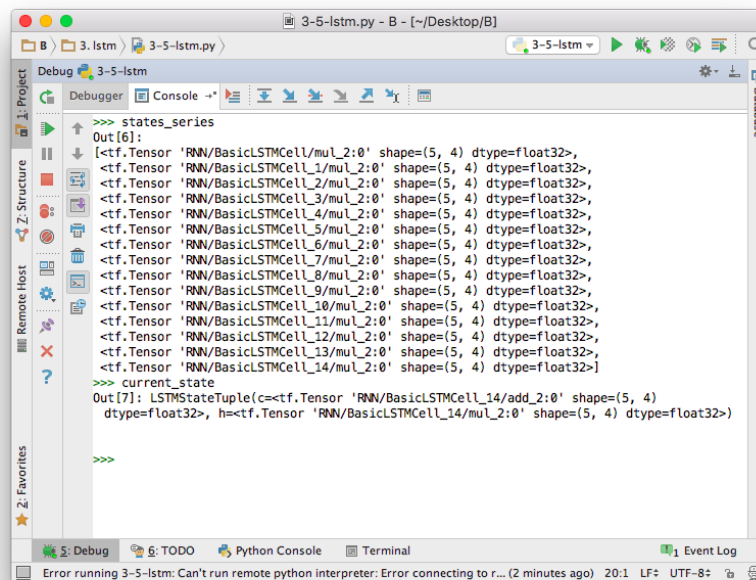
the second is the hidden state. So you need to change line 28 where the `init_state` is placeholders are declared to these lines:

```
1 cell_state = tf.placeholder(tf.float32, [batch_size, st
2 hidden_state = tf.placeholder(tf.float32, [batch_size,
3 init_state = tf.nn.rnn_cell.LSTMStateTuple(cell_state,
```

Changing the forward pass is now straight forward, you just change the function call to create a LSTM and supply the initial state-tuple on line 38–39.

```
1 cell = tf.nn.rnn_cell.BasicLSTMCell(state_size, state_i
2 states_series, current_state = tf.nn.rnn(cell, inputs_s
```

The `states_series` will be a list of *hidden states* as tensors, and `current_state` will be a `LSTMStateTuple` which shows both the *hidden-* and the *cell state* on the last time-step as shown below:



Outputs of the previous states and the last `LSTMStateTuple`

So the `current_state` returns the cell- and hidden state in a tuple. They should be separated after calculation and supplied to the

placeholders in the run-function on line 90.

```
1  _total_loss, _train_step, _current_state, _predictions
2      [total_loss, train_step, current_state, prediction
3      feed_dict={
4          batchX_placeholder: batchX,
5          batchY_placeholder: batchY,
6          cell_state: _current_cell_state,
7          hidden_state: _current_hidden_state
8      ..
```

Whole program

This is the full code for creating a RNN with Long short-term memory.

```

1  from __future__ import print_function, division
2  import numpy as np
3  import tensorflow as tf
4  import matplotlib.pyplot as plt
5
6  num_epochs = 100
7  total_series_length = 50000
8  truncated_backprop_length = 15
9  state_size = 4
10 num_classes = 2
11 echo_step = 3
12 batch_size = 5
13 num_batches = total_series_length//batch_size//trunca
14
15 def generateData():
16     x = np.array(np.random.choice(2, total_series_len
17     y = np.roll(x, echo_step)
18     y[0:echo_step] = 0
19
20     x = x.reshape((batch_size, -1)) # The first inde
21     y = y.reshape((batch_size, -1))
22
23     return (x, y)
24
25 batchX_placeholder = tf.placeholder(tf.float32, [batch
26 batchY_placeholder = tf.placeholder(tf.int32, [batch_
27
28 cell_state = tf.placeholder(tf.float32, [batch_size,
29 hidden_state = tf.placeholder(tf.float32, [batch_size
30 init_state = tf.nn.rnn_cell.LSTMStateTuple(cell_state
31
32 W2 = tf.Variable(np.random.rand(state_size, num_class
33 b2 = tf.Variable(np.zeros((1,num_classes)), dtype=tf.
34
35 # Unpack columns
36 inputs_series = tf.split(1, truncated_backprop_length
37 labels_series = tf.unpack(batchY_placeholder, axis=1)
38
39 # Forward passes
40 cell = tf.nn.rnn_cell.BasicLSTMCell(state_size, state
41 states_series, current_state = tf.nn.rnn(cell, inputs

```

```

42
43 logits_series = [tf.matmul(state, W2) + b2 for state
44 predictions_series = [tf.nn.softmax(logits) for logit
45
46 losses = [tf.nn.sparse_softmax_cross_entropy_with_log
47 total_loss = tf.reduce_mean(losses)
48
49 train_step = tf.train.AdagradOptimizer(0.3).minimize(
50
51 def plot(loss_list, predictions_series, batchX, batch
52     plt.subplot(2, 3, 1)
53     plt.cla()
54     plt.plot(loss_list)
55
56     for batch_series_idx in range(5):
57         one_hot_output_series = np.array(predictions_
58         single_output_series = np.array([(1 if out[0]
59
60         plt.subplot(2, 3, batch_series_idx + 2)
61         plt.cla()
62         plt.axis([0, truncated_backprop_length, 0, 2]
63         left_offset = range(truncated_backprop_length
64         plt.bar(left_offset, batchX[batch_series_idx,
65         plt.bar(left_offset, batchY[batch_series_idx,
66         plt.bar(left_offset, single_output_series * 0
67
68     plt.draw()
69     plt.pause(0.0001)
70

```

Next step

In the next article we will create a multi-layered or “deep” recurrent neural network, also with long short-term memory.

