



Erik Hallström [Follow](#)

Studied Engineering Physics and in Machine Learning at Royal Institute of Technology in Stockholm. Also...

Nov 19, 2016 · 3 min read

Using the Multilayered LSTM API in TensorFlow (4/7)

In the [previous article](#) we learned how to use the TensorFlow API to create a Recurrent neural network with Long short-term memory. In this post we will make that architecture deep, introducing a LSTM with multiple layers.

One thing to notice is that for every layer of the network we will need a hidden state and a cell state. Typically the input to the next LSTM-layer will be the previous state for that particular layer as well as the hidden activations of the “lower” or previous layer. There is a good diagram in [this article](#).

We could continue to store the states for each layer in many `LSTMTuples`, but that would require a lot of overhead. You can only input data to the placeholders through the `feed_dict` as Python lists or Numpy arrays anyways (not as `LSTMTuples`) so we still would have to convert between the datatypes. Why not save the whole state for the network in a big tensor? In order to do this the first thing we want to do is to replace `_current_cell_state` and `_current_hidden_state` on line 81–82 with the more generic:

```
1  _current_state = np.zeros((num_layers, 2, batch_size, s
2
```

You also have to declare the new setting `num_layers = 3` in the beginning of the file, but you may choose any number of layers. The “2” refers to the two states, cell- and hidden-state. So for each layer and each sample in a batch, we have both a cell state and a hidden state vector with the size `state_size`.

Now modify lines 93 to 103 (the run function and the separation of the state tuple) back to the original statement, since the state is now stored in a single tensor.

```

1  _total_loss, _train_step, _current_state, _predictions_
2      [total_loss, train_step, current_state, predictions
3      feed_dict={
4          batchX_placeholder: batchX,
5          batchY_placeholder: batchY,
6          init_state: _current_state

```

You can change these lines 28 to 30 in the previous post:

```

1  cell_state = tf.placeholder(tf.float32, [batch_size, st
2  hidden_state = tf.placeholder(tf.float32, [batch_size,
3  init_state = tf.nn.rnn_cell.LSTMStateTuple(cell_state,

```

To a single placeholder containing the whole state.

```

1  init_state = tf.placeholder(tf.float32, [num_layers, 2,
2

```

Since the TensorFlow Multilayer-LSTM-API accepts the state as a tuple of LSTM tuples, we need to unpack the state state into this structure. For each layer in the state we then create a `LSTMStateTuple` stated, and put these in a tuple, as shown below. Add this just after the `init_state` placeholder.

```

1  state_per_layer_list = tf.unpack(init_state, axis=0)
2  rnn_tuple_state = tuple(
3      [tf.nn.rnn_cell.LSTMStateTuple(state_per_layer_list
4      for idx in range(num_layers)]
5  )

```

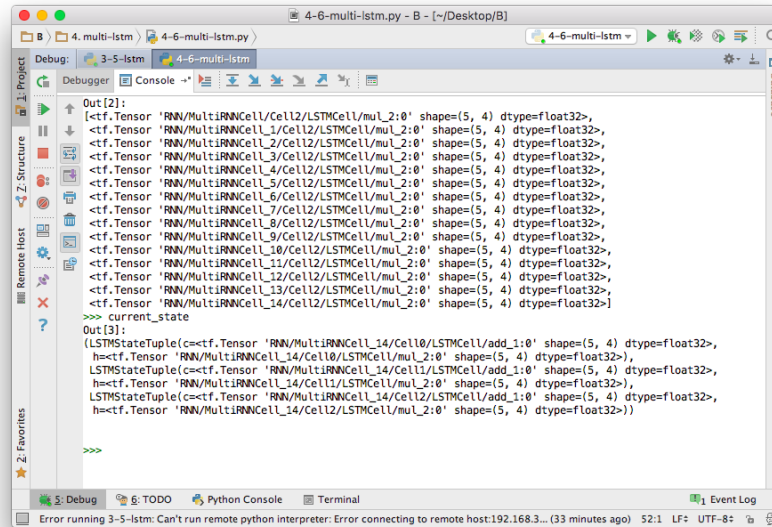
The forward pass on lines 40 and 41 should be changed to this:

```

1  # Forward passes
2  cell = tf.nn.rnn_cell.LSTMCell(state_size, state_is_tup
3  cell = tf.nn.rnn_cell.MultiRNNCell([cell] * num_layers,
4  states_series, current_state = tf.nn.rnn(cell, inputs_s

```

The multi-layered LSTM is created by first making a single `LSMTCell` , and then duplicating this cell in an array, supplying it to the `MultiRNNCell` API call. The forward pass uses the usual `tf.nn.rnn` , let's print the output of this function, the `states_series` and `current_state` variables.



```

Out[2]:
[<tf.Tensor 'RNN/MultiRNNCell_12/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_11/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_10/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_9/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_8/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_7/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_6/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_5/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_4/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_3/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_2/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_1/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_14/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_13/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_12/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_11/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_10/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_9/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_8/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_7/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_6/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_5/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_4/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_3/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_2/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>,
 <tf.Tensor 'RNN/MultiRNNCell_1/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>]
>>> current_state
Out[3]:
(LSTMStateTuple(c=<tf.Tensor 'RNN/MultiRNNCell_14/Cell0/LSTMCell/add_1:0' shape=(5, 4) dtype=float32>,
 h=<tf.Tensor 'RNN/MultiRNNCell_14/Cell0/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>),
 LSTMStateTuple(c=<tf.Tensor 'RNN/MultiRNNCell_14/Cell1/LSTMCell/add_1:0' shape=(5, 4) dtype=float32>,
 h=<tf.Tensor 'RNN/MultiRNNCell_14/Cell1/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>),
 LSTMStateTuple(c=<tf.Tensor 'RNN/MultiRNNCell_14/Cell2/LSTMCell/add_1:0' shape=(5, 4) dtype=float32>,
 h=<tf.Tensor 'RNN/MultiRNNCell_14/Cell2/LSTMCell/mul_2:0' shape=(5, 4) dtype=float32>))
>>>

```

Output of the previous states and the last LSTMStateTuples

Take a look at the tensor names between single quotes, we see that the RNN is unrolled 15 times. In the `states_series` all outputs have the name “Cell2”, it means that we get the output of the last LSTM layer’s hidden state in the list. Furthermore the `LSTMStateTuple` in the `current_state` gives the whole state of all layers in the network. “Cello” refers to the first layer, “Cell1” to the second and “Cell2” to the third and final layer, “h” and “c” refers to hidden- and cell state.

Whole program

This is the whole self-contained script, just copy and run.

```

1  from __future__ import print_function, division
2  import numpy as np
3  import tensorflow as tf
4  import matplotlib.pyplot as plt
5
6  num_epochs = 100
7  total_series_length = 50000
8  truncated_backprop_length = 15
9  state_size = 4
10 num_classes = 2
11 echo_step = 3
12 batch_size = 5
13 num_batches = total_series_length//batch_size//truncated_backprop_length
14 num_layers = 3
15
16 def generateData():
17     x = np.array(np.random.choice(2, total_series_length, dtype=int))
18     y = np.roll(x, echo_step)
19     y[0:echo_step] = 0
20
21     x = x.reshape((batch_size, -1)) # The first index is the batch
22     y = y.reshape((batch_size, -1))
23
24     return (x, y)
25
26 batchX_placeholder = tf.placeholder(tf.float32, [batch_size, -1])
27 batchY_placeholder = tf.placeholder(tf.int32, [batch_size, -1])
28
29 init_state = tf.placeholder(tf.float32, [num_layers, state_size])
30
31 state_per_layer_list = tf.unpack(init_state, axis=0)
32 rnn_tuple_state = tuple(
33     [tf.nn.rnn_cell.LSTMStateTuple(state_per_layer_list[idx],
34                                     tf.zeros(state_size, dtype=tf.float32))
35      for idx in range(num_layers)]
36 )
37
38 w2 = tf.Variable(np.random.rand(state_size, num_classes), dtype=tf.float32)
39 b2 = tf.Variable(np.zeros((1, num_classes)), dtype=tf.float32)
40
41 # Unpack columns
42 inputs_series = tf.split(1, truncated_backprop_length, batchX_placeholder)
43 outputs_series = tf.zeros([truncated_backprop_length, num_classes], dtype=tf.float32)

```

```

42 labels_series = tf.unpack(batchY_placeholder, axis=1)
43
44 # Forward passes
45 cell = tf.nn.rnn_cell.LSTMCell(state_size, state_is_t
46 cell = tf.nn.rnn_cell.MultiRNNCell([cell] * num_layer
47 states_series, current_state = tf.nn.rnn(cell, inputs
48
49 logits_series = [tf.matmul(state, W2) + b2 for state
50 predictions_series = [tf.nn.softmax(logits) for logit
51
52 losses = [tf.nn.sparse_softmax_cross_entropy_with_log
53 total_loss = tf.reduce_mean(losses)
54
55 train_step = tf.train.AdagradOptimizer(0.3).minimize(
56
57 def plot(loss_list, predictions_series, batchX, batch
58     plt.subplot(2, 3, 1)
59     plt.cla()
60     plt.plot(loss_list)
61
62     for batch_series_idx in range(5):
63         one_hot_output_series = np.array(predictions_
64         single_output_series = np.array([(1 if out[0]
65
66         plt.subplot(2, 3, batch_series_idx + 2)
67         plt.cla()
68         plt.axis([0, truncated_backprop_length, 0, 2]
69         left_offset = range(truncated_backprop_length
70         plt.bar(left_offset, batchX[batch_series_idx,
71         plt.bar(left_offset, batchY[batch_series_idx,
72         plt.bar(left_offset, single_output_series * 0

```

Next step

In the next [article](#) we will speed up the graph creation by not splitting up our inputs and labels into a Python list.

